

Process management

Examples and exercises

EPL222 – Lab2

Summary of fork()

- ▶ The syntax of the fork() system function is as follows:

```
pid_t fork(void);
```

- ▶ The fork() system function does not accept any argument
 - It returns an integer of the type `pid_t` (defined in library `<sys/types.h>`)
- ▶ On success, fork() returns the PID of the child process which is greater than 0
 - Inside the child process, the return value is 0
 - If fork() fails, then it returns -1



Simple pipes (between parent and child)

- ▶ **pipe()** is used for passing information from one process to another
 - **pipe()** is unidirectional therefore, for two-way communication between processes, two pipes can be set up, one for each direction. E.g.:

- ```
int fd[2];
pipe(fd);
fd[0]; //-> for using read end
fd[1]; //-> for using write end
```



# Pipe example

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
int main() {
 int fd[2], i = 0;
 pipe(fd);
 pid_t pid = fork();

 if(pid > 0) {
 wait(NULL);
 close(fd[1]);

 int arr[10];
 int n = read(fd[0], arr, sizeof(arr));

 for (i = 0; i < n/4; i++) {
 printf("%d ", arr[i]);
 }
 } else if(pid == 0) {
 int arr[] = {1, 2, 3, 4, 5};

 close(fd[0]);
 write(fd[1], arr, sizeof(arr));
 } else {
 perror("error\n"); //fork()
 }
}
```

// wait for child to finish  
// no need to use the write end of pipe here so close it

// n stores the total bytes read successfully

// printing the array received from child process

// no need to use the read end of pipe here so close it



# Wait system calls

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

- ▶ All these system calls are used to wait for state changes in a child of the calling process and obtain information about the child whose state has changed. A state change is considered to be:
  - the child terminated;
  - the child was stopped by a signal;
  - or the child was resumed by a signal.
- ▶ In the case of a terminated child, performing a wait allows the system to release the resources associated with the child; if a wait is not performed, then the terminated child remains in a "zombie" state
- ▶ If a child has already changed state, then these calls return immediately. Otherwise, they block until either a child changes state or a signal handler interrupts the call



# Wait system calls

- ▶ The **wait()** system call suspends execution of the current process until one of its children terminates.
  - The call **wait(&status)** is equivalent to:  
`waitpid(-1, &status, 0);`
- ▶ The **waitpid()** system call suspends execution of the current process until a child specified by **pid argument** has changed state
  - By default, **waitpid()** waits only for terminated children, but this behaviour is modifiable via the options argument
- ▶ The **waitid()** system call (available since Linux 2.6.9) provides more precise control over which child state changes to wait for



# Wait system calls

- ▶ If only one child process is terminated, `wait()` returns the process ID of the terminated child
- ▶ If more than one child processes are terminated, then `wait()` catches any *arbitrarily child* and returns the process ID of that child
- ▶ If a process has no child process, then `wait()` immediately returns “-1”
- ▶ When `wait()` returns it also defines an *exit status* (which tells us why the process terminated)
  - Done via a pointer passed in the call of `wait()`



# Wait example

```
#include<stdio.h>
#include<stdlib.h>
#include<sys/wait.h>
#include<sys/types.h>

int main() {
 pid_t cpid;
 if (fork() == 0)
 exit(0);
 else
 cpid = wait(NULL);

 printf("Parent pid = %d\n", getpid());
 printf("Child pid = %d\n", cpid);
 return 0;
}
```





# Checking the Status of child

- ▶ To learn about the exit status of a program we can use the macros from `<sys/wait.h>` which check the termination status and return the exit status
  - **WIFEXITED(status)** returns true if the child terminated normally, that is, by calling `exit()` or `_exit()`, or by returning from `main()`
  - **WEXITSTATUS(status)** returns the exit status of the child. This consists of the least significant 8 bits of the status argument that the child specified in a call to `exit()` or `_exit()` or as the argument for a return statement in `main()`
    - This macro should only be employed if **WIFEXITED** returned true



# Checking the Status of child

## Example

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/wait.h>

int main(){
 pid_t c_pid, pid;
 int status;
 c_pid = fork();

 if(c_pid == 0){ //child
 pid = getpid();
 printf("Child: %d: I'm the child\n", pid, c_pid);
 printf("Child: sleeping for 2-seconds, then exiting with status 12\n");

 sleep(2); //sleep for 2 seconds
 exit(12); //exit with status 12
 } else if (c_pid > 0){ //parent
 pid = wait(&status); //waiting for child to terminate
 if (WIFEXITED(status)){
 printf("Parent: Child exited with status: %d\n", WEXITSTATUS(status));
 }
 } else { //error: The return of fork() is negative
 perror("fork failed");
 exit(2); //exit failure, hard
 }
 return 0;
}
```



# waitpid() options

The value of *options* is the bitwise OR of zero or more of the following constants:

| Tag        | Description                                                                                                                                                                     |
|------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WNOHANG    | return immediately if no child has exited.                                                                                                                                      |
| WUNTRACED  | also return if a child has stopped (but not traced via <code>ptrace()</code> ). Status for traced children which have stopped is provided even if this option is not specified. |
| WCONTINUED | (Since Linux 2.6.10) also return if a stopped child has been resumed by delivery of SIGCONT.                                                                                    |



# wait() and waitpid() Status Information

| Tag                  | Description                                                                                                                                                                                                                                                                                                                                    |
|----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WIFEXITED(status)    | returns true if the child terminated normally, that is, by calling <code>exit()</code> or <code>_exit()</code> , or by returning from <code>main()</code> .                                                                                                                                                                                    |
| WEXITSTATUS(status)  | returns the exit status of the child. This consists of the least significant 16–8 bits of the status argument that the child specified in a call to <code>exit()</code> or <code>_exit()</code> or as the argument for a return statement in <code>main()</code> . This macro should only be employed if <code>WIFEXITED</code> returned true. |
| WIFSIGNALED(status)  | returns true if the child process was terminated by a signal.                                                                                                                                                                                                                                                                                  |
| WTERMSIG(status)     | returns the number of the signal that caused the child process to terminate. This macro should only be employed if <code>WIFSIGNALED</code> returned true.                                                                                                                                                                                     |
| WCOREDUMP(status)    | returns true if the child produced a core dump. This macro should only be employed if <code>WIFSIGNALED</code> returned true. This macro is not specified in POSIX.1–2001 and is not available on some Unix implementations (e.g., AIX, SunOS). Only use this enclosed in <code>#ifdef WCOREDUMP ... #endif</code> .                           |
| WIFSTOPPED(status)   | returns true if the child process was stopped by delivery of a signal; this is only possible if the call was done using <code>WUNTRACED</code> or when the child is being traced.                                                                                                                                                              |
| WSTOPSIG(status)     | returns the number of the signal which caused the child to stop. This macro should only be employed if <code>WIFSTOPPED</code> returned true.                                                                                                                                                                                                  |
| WIFCONTINUED(status) | (Since Linux 2.6.10) returns true if the child process was resumed by delivery of <code>SIGCONT</code> .                                                                                                                                                                                                                                       |



# Exercise 1

- ▶ Write a program to create 4 processes: parent process and its child process which perform various tasks on an array of 10 numbers:
  - Parent process count the frequency of a number
  - 1<sup>st</sup> child sort the array
  - 2<sup>nd</sup> child find total even number(s) in the array
  - 3<sup>rd</sup> child calculate the sum of even numbers in the array



# Exercise 2

- ▶ Write a C program, fibchild.c, which computes and prints Fibonacci numbers in child processes
  - The program runs in a loop and creates a child process, which computes the current Fibonacci number and prints it.
  - The parent waits for the child finish and then continues executing.
  - Take the number of Fibonacci numbers to be printed as input.
- ▶ Fibonacci number are defined as:
  - $fib(n) = fib(n-1) + fib(n-2), n > 2$
  - $fib(1) = fib(2) = 1$

