

Threads in Java

Implementing Monitors

EPL222 – Lab8

Producer – Consumer

```
class Producer implements Runnable {  
    private final Queue<Integer> queue;  
    private final int maxSize;  
    public Producer(Queue<Integer> queue, int maxSize) {  
        this.queue = queue; this.maxSize = maxSize;  
    }  
    public void run() {  
        int i = 0;  
        while (true) {  
            synchronized (queue) {  
                while (queue.size() == maxSize) {  
                    try {  
                        System.out.println(Thread.currentThread().getName() +  
                            ": Buffer is full");  
                        queue.wait();  
                    } catch (InterruptedException e) {  
                        e.printStackTrace();  
                    }  
                }  
                i++;  
                System.out.println(Thread.currentThread().getName() +  
                    ": Producing value: " + i);  
                queue.add(i);  
                queue.notifyAll();  
            }  
        }  
    }  
}
```

Producer

```
class Consumer implements Runnable {  
    private final Queue<Integer> queue;  
  
    public Consumer(Queue<Integer> queue) {  
        this.queue = queue;  
    }  
  
    public void run() {  
        int i;  
        while (true) {  
            synchronized (queue) {  
                while (queue.isEmpty()) {  
                    try {  
                        queue.wait();  
                    } catch (InterruptedException e) {  
                        e.printStackTrace();  
                    }  
                }  
                i = queue.remove();  
                System.out.println(Thread.currentThread().  
                    getName() + ": Consuming value: " + i);  
                queue.notifyAll();  
            }  
        }  
    }  
}
```

Consumer

Producer – Consumer (continued)

```
public class ProducerConsumerExample {  
    public static void main(String[] args) {  
        final Queue<Integer> buffer = new LinkedList<>();  
        int maxsize = 100, producers = 3, consumers = 5;  
        Thread t[] = new Thread[producers+consumers];  
  
        for (int i = 0; i < producers; i++) {  
            t[i] = new Thread(new Producer(buffer, maxsize), "Producer " + i);  
        }  
        for (int i = 0; i < consumers ; i++) {  
            t[producers+i] = new Thread(new Consumer(buffer), "Consumer " + i);  
        }  
        for (int i = 0; i < t.length; i++) {  
            t[i].start();  
        }  
    }  
}
```



Problems with previous solution

- ▶ This is **NOT a monitor!**
 - Two separate classes that implement the producer and consumer threads
 - The used **resource (the *queue*) is not protected** from the threads
 - instead the threads have direct access to them!
 - Threads take care of their concurrency!
- ▶ Does it work?
 - The provided solution is correct as far as concurrency is concerned
 - This is because all producers and consumers threads **synchronize on the same object** (there is only one *queue* object) but each thread executes on a different object!
 - Each producer thread continues producing from the last value that thread produced!
- ▶ We have no way to wake a producer or a consumer!
 - Both consumers and producers wait on the same object!
 - We have to use `queue.notifyAll()`



Producer – Consumer (continued)

```
public class ProducerConsumerExample2 {  
    public static void main(String[] args) {  
        final Queue<Integer> buffer = new LinkedList<>();  
        int maxsize = 10, producers = 3, consumers = 5;  
        Thread t[] = new Thread[producers + consumers];
```

**Incorrect
concurrency!**

Only 1
producer
and 1
consumer
object!

```
        Producer p = new Producer(buffer, maxsize);  
        Consumer c = new Consumer(buffer);
```

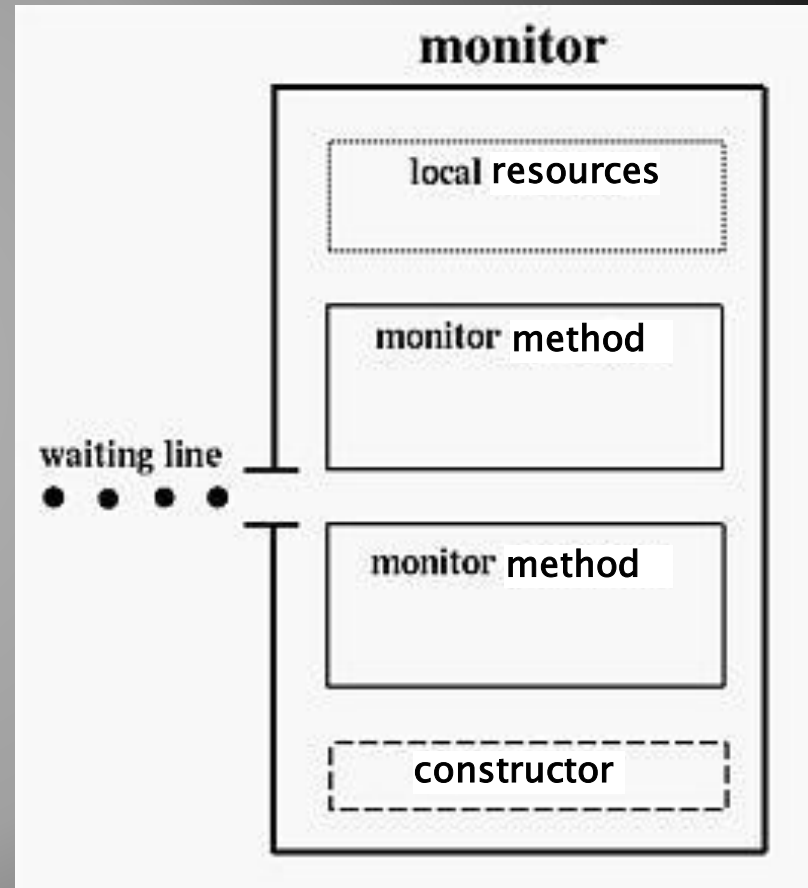
```
        for (int i = 0; i < producers; i++) {  
            t[i] = new Thread(p, "Producer " + i);  
        }  
        for (int i = 0; i < consumers; i++) {  
            t[producers + i] = new Thread(c, "Consumer " + i);  
        }  
        for (int i = 0; i < t.length; i++) {  
            t[i].start();  
        }  
    }  
}
```

All producer threads
share the same variable
to produce values!



Implementing a Monitor

- ▶ Remember that a monitor has some **local resources**, monitor **procedures (methods)** and an **initialization** part (**constructor**)
- ▶ In Java a monitor can be a class
 - The public methods of the class form the boundary of the monitor
 - Making them synchronized guarantees that, at any time, there can only be **one thread executing within the monitor** boundary



Implementing a Monitor in Java

- ▶ To implement a monitor in Java, we need to address the same three questions as in C:
 - (1) how do we make sure that the local variables/objects will not be accessed by non-monitor procedures
 - (2) how do we make sure that the user can only “see” the interface without knowing the details of the monitor
 - (3) how do we properly setup the monitor boundary so that mutual exclusion can be guaranteed
- ▶ Questions (1) and (2) have a simple solution
 - Just make everything beyond the monitor methods private!
 - Or define an interface with just the monitor methods
 - The implantation class can have more (private) methods and/or local variables
- ▶ Question (3) has a natural answer: make the monitor methods synchronized!



Producer Consumer Monitor – 1st attempt

```
public class Producer implements Runnable {
    private final PCMonitor m;

    public Producer(PCMonitor m) {
        this.m = m;
    }

    public void run() {
        while (true) {
            System.out.println(Thread.currentThread()
                               .getName() + ": Producing value: " +
                               m.produce());

            try {
                if (Math.random() > 0.5)
                    Thread.sleep((int) (Math.random()*1000)%1000);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```

Producer

```
public class Consumer implements Runnable {
    private final PCMonitor m;

    public Consumer(PCMonitor m) {
        this.m = m;
    }

    public void run() {
        while (true) {
            System.out.println(Thread.currentThread()
                               .getName() + ": Consuming value: " +
                               m.consume());

            try {
                if (Math.random() > 0.5)
                    Thread.sleep((int) (Math.random()*1000)%1000);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```

Consumer

Producer Consumer Monitor – 1st attempt

```
public interface PCMonitor {  
  
    int produce();  
    int consume();  
  
}
```

```
import java.util.LinkedList;  
import java.util.Queue;
```

```
public class PCMonitorImpl  
    implements PCMonitor {  
    private final Queue<Integer> queue;  
    private final int maxSize;  
    private int value = 0;  
  
    public PCMonitorImpl(int maxSize) {  
        this.queue = new LinkedList<>();  
        this.maxSize = maxSize;  
    }  
}
```

```
    public synchronized int produce() {  
        while (queue.size() == maxSize) {  
            try {  
                System.out.println(Thread.currentThread().getName()+" : Buffer full");  
                wait();  
            } catch (InterruptedException e) {  
                e.printStackTrace();  
            }  
        }  
        queue.add(++value);  
        notifyAll();  
        return value;  
    }  
  
    public synchronized int consume() {  
        int i;  
        while (queue.isEmpty()) {  
            try {  
                wait();  
            } catch (InterruptedException e) {  
                e.printStackTrace();  
            }  
        }  
        i = queue.remove();  
        notifyAll();  
        return i;  
    }  
}
```

Monitor



Producer Consumer Monitor – 1st attempt

```
public class PCMExample {  
    public static void main(String[] args) {  
        int maxSize = 10, producers = 3, consumers = 5;  
  
        PCMonitor m = new PCMonitorImpl(maxSize);  
        Thread t[] = new Thread[producers + consumers];  
  
        for (int i = 0; i < producers; i++) {  
            t[i] = new Thread(new Producer(m), "Producer " + i);  
        }  
        for (int i = 0; i < consumers; i++) {  
            t[producers + i] = new Thread(new Consumer(m), "Consumer " + i);  
        }  
        for (int i = 0; i < t.length; i++) {  
            t[i].start();  
        }  
    }  
}
```



Problems with 1st monitor attempt

- ▶ Now it is **a monitor!**
 - The classes that implement the producer and consumer threads are separate from the monitor
 - The used **resource (the queue) NOT accessible** from the threads
 - The monitor class takes care of concurrency!
- ▶ We have no way to wake a producer or a consumer!
 - Both consumers and producers wait on the same object – **the monitor object!**
 - We must use `notifyAll()`



Condition variables in Java

Interface Condition

- ▶ **Condition** factors out the Object monitor methods (**wait**, **notify** and **notifyAll**) into distinct objects to give the effect of having multiple wait-sets per object, by combining them with the use of arbitrary **Lock** implementations
 - A **Lock** replaces the use of synchronized methods and statements
 - A **Condition** replaces the use of the Object monitor methods.
 - It is Java's equivalent to *condition variables*
 - As always access to a condition must be protected, so a lock of some form must be associated with it
 - The key property that waiting for a condition provides is that it *atomically* releases the associated lock and suspends the current thread, just like **Object.wait**
- ▶ A **Condition** instance is intrinsically bound to a lock
 - To obtain a **Condition** instance for a particular **Lock** instance use its **newCondition()** method



The Lock interface

- ▶ **Lock** implementations provide more extensive locking operations than can be obtained using synchronized methods and statements
 - They allow more flexible structuring, may have quite different properties, and may support multiple associated **Condition** objects
 - Only one thread at a time can acquire a given *lock instance*
- ▶ The use of synchronized methods or statements provides access to the implicit monitor lock associated with every object, but forces all lock acquisition and release to occur in a block-structured way
 - If multiple locks are acquired they must be released in the opposite order
 - All locks must be released in the **same lexical scope in which they were acquired**
 - There are occasions where you need to work with locks in a more flexible way
 - E.g., the use of "hand-over-hand" or "chain locking": you acquire the lock of node A, then node B, then release A and acquire C, then release B and acquire D and so on
 - Implementations of the Lock interface enable the use of such techniques by allowing a lock to be acquired and released **in different scopes**, and allowing multiple locks to be **acquired and released in any order**



The Lock interface

- ▶ With this increased flexibility comes additional responsibility
 - The absence of block-structured locking removes the automatic release of locks that occurs with synchronized methods and statements – the following idiom should be used:

```
Lock l = ...;  
l.lock();  
try {  
    // access the resource protected by this lock  
} finally {  
    l.unlock();  
}
```

- When locking and unlocking occur in different scopes, care must be taken to ensure that all code that is executed while the lock is held is protected by try-finally or try-catch to ensure that the lock is released when necessary.
 - Lock implementations provide additional functionality over the use of synchronized methods and statements by providing a non-blocking attempt to acquire a lock (`tryLock()`, `tryLock(long, TimeUnit)`)
- ▶ Note that Lock instances are just normal objects and can themselves be used as the target in a synchronized statement
 - Acquiring the monitor lock of a Lock instance has no specified relationship with invoking any of the `lock()` methods of that instance
 - It is recommended that to avoid confusion you **never use Lock instances in this way**



Condition and Lock Interfaces

Type	Method and Description
void	<code>await()</code> Causes the current thread to wait until it is signalled or interrupted.
boolean	<code>await(long time, TimeUnit unit)</code> Causes the current thread to wait until it is signalled or interrupted, or the specified waiting time elapses.
long	<code>awaitNanos(long nanosTimeout)</code> Causes the current thread to wait until it is signalled or interrupted, or the specified waiting time elapses.
void	<code>awaitUninterruptibly()</code> Causes the current thread to wait until it is signalled.
boolean	<code>awaitUntil(Date deadline)</code> Causes the current thread to wait until it is signalled or interrupted, or the specified deadline elapses.
void	<code>signal()</code> Wakes up one waiting thread.
void	<code>signalAll()</code> Wakes up all waiting threads.

Interface Condition

Type	Method and Description
void	<code>lock()</code> Acquires the lock.
void	<code>lockInterruptibly()</code> Acquires the lock unless the current thread is interrupted.
Condition	<code>newCondition()</code> Returns a new Condition instance that is bound to this Lock instance.
boolean	<code>tryLock()</code> Acquires the lock only if it is free at the time of invocation.
boolean	<code>tryLock(long time, TimeUnit unit)</code> Acquires the lock if it is free within the given waiting time and the current thread has not been interrupted.
void	<code>unlock()</code> Releases the lock.

All Known Implementing Classes:

`ReentrantLock`, `ReentrantReadWriteLock.ReadLock`,
`ReentrantReadWriteLock.WriteLock`

Interface Lock

Producer Consumer Monitor – 2nd attempt

```
import java.util.LinkedList;
import java.util.Queue;
import java.util.concurrent.locks.Condition;
import java.util.concurrent.locks.Lock;
import java.util.concurrent.locks.ReentrantLock;

public class PCMonitorImp2 implements PCMonitor {
    private final Lock lock = new ReentrantLock();
    private final Condition qFull = lock.newCondition();
    private final Condition qEmpty = lock.newCondition();

    private final Queue<Integer> queue;
    private final int maxSize;
    private int value = 0;

    public PCMonitorImp2(int maxSize) {
        this.queue = new LinkedList<>();
        this.maxSize = maxSize;
    }
}
```

```
public interface PCMonitor {
    int produce();
    int consume();
}
```

Monitor



Producer Consumer Monitor – 2nd attempt

```
public int produce() {  
    lock.lock();  
    try {  
        while (queue.size() == maxSize) {  
            System.out.println(Thread.currentThread()  
                .getName() + ": Buffer is full");  
            qFull.await();  
        }  
        queue.add(++value);  
        qEmpty.signal();  
        return value;  
    } catch (InterruptedException e) {  
        e.printStackTrace();  
        return -1;  
    } finally {  
        lock.unlock();  
    }  
}
```

```
public int consume() {  
    lock.lock();  
    int i = 0;  
    try {  
        while (queue.isEmpty())  
            qEmpty.await();  
        i = queue.remove();  
        qFull.signal();  
        return i;  
    } catch (InterruptedException e) {  
        e.printStackTrace();  
        return -1;  
    } finally {  
        lock.unlock();  
    }  
}
```

